Deep Reinforcement Learning

Beril Sirmacek

7th March 2019, Enschede

Outline

- Reinforcement Learning (RL)
- Deep reinforcement learning (DRL)
- Practical applications of RL

What is reinforcement learning?

- Similarity with other deep learning nets and RL

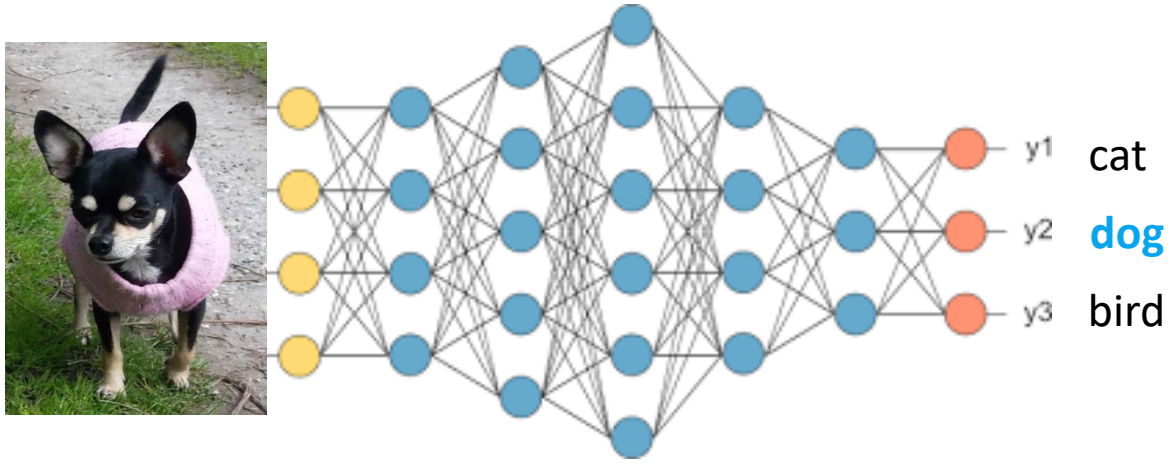

Supervised learning

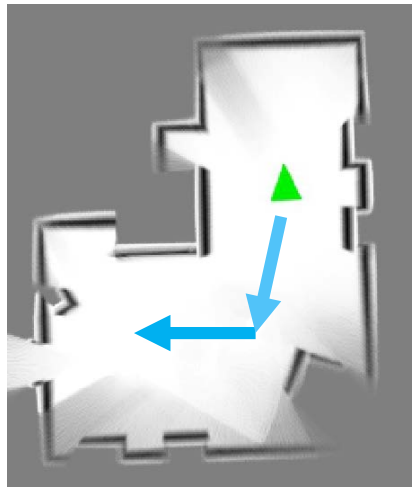

Reinforcement learning

decision making!

What is reinforcement learning?

- Difference from other deep learning nets and RL

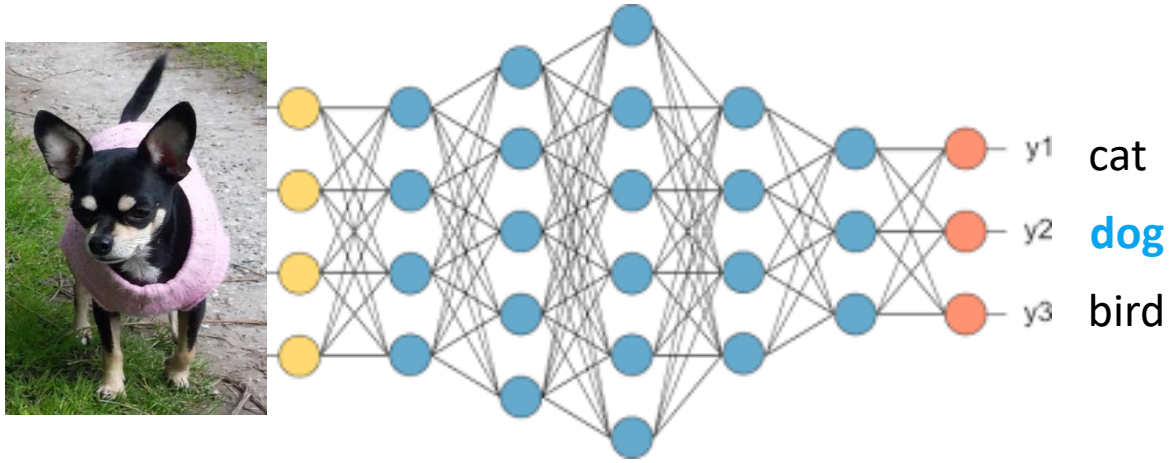

Supervised learning

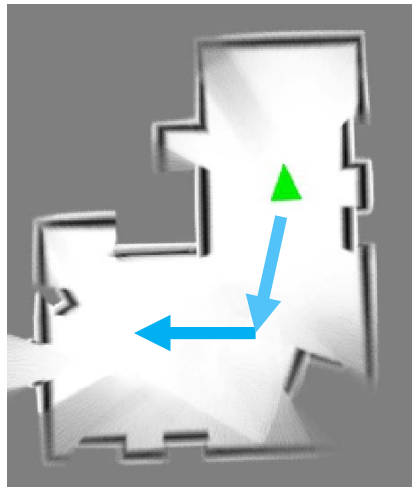

Reinforcement learning

A chain of decisions (actions) & dependency on the previous decisions

What is reinforcement learning?

- Difference from other deep learning nets and RL

Where is UNSUPERVISED LEARNING?

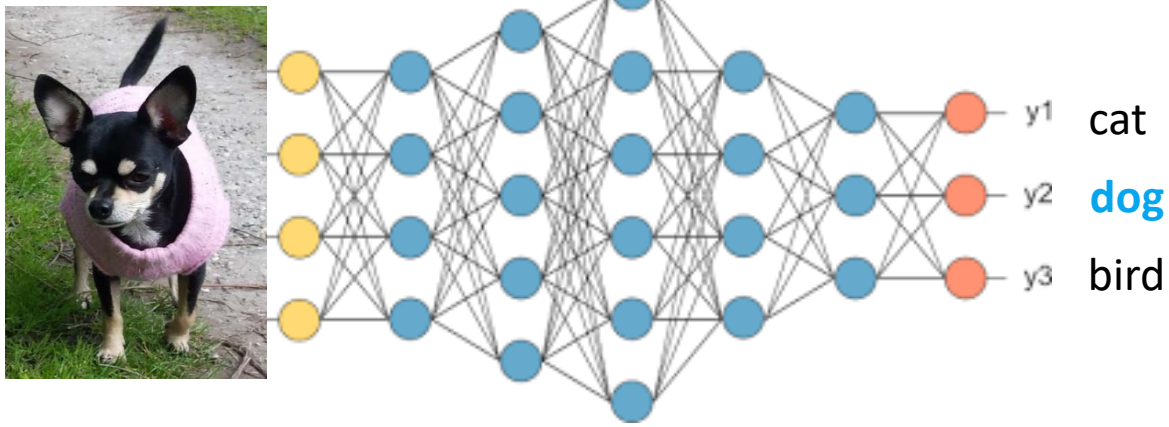

Supervised learning

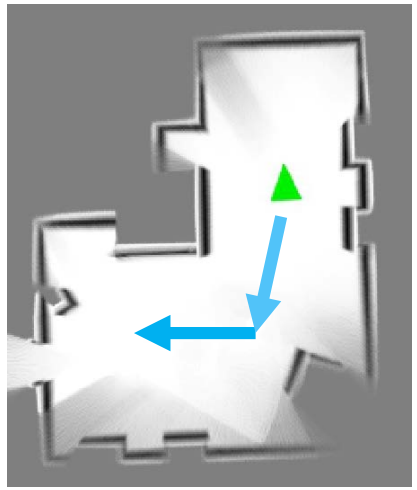

Reinforcement learning

A chain of decisions (actions) & dependency on the previous decisions

What is reinforcement learning?

- Here are the examples / Loss function
- Learn it (learning by example)

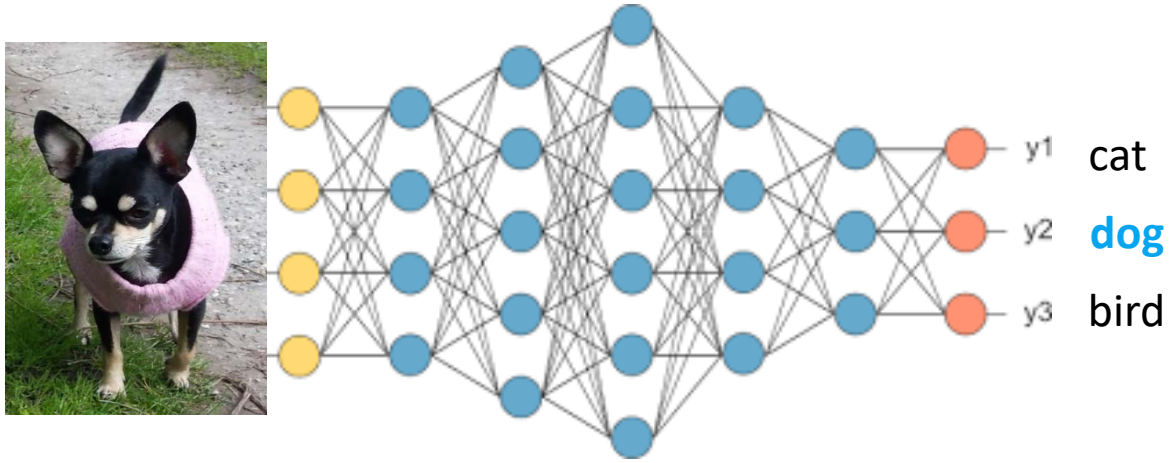

Supervised learning

- Here is the environment
- Learn it (learning by experience)

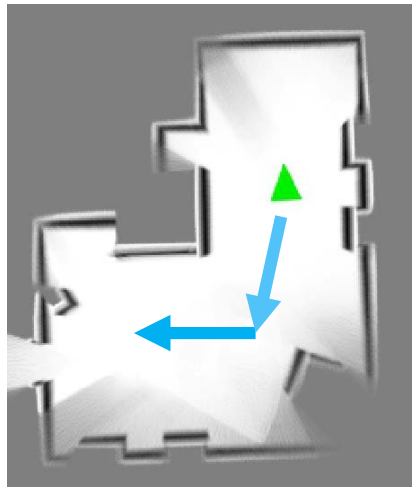

Reinforcement learning

Problem description

- What series of actions will give me the highest reward?
  (or helps me to reach to my goal faster/safer/better?)

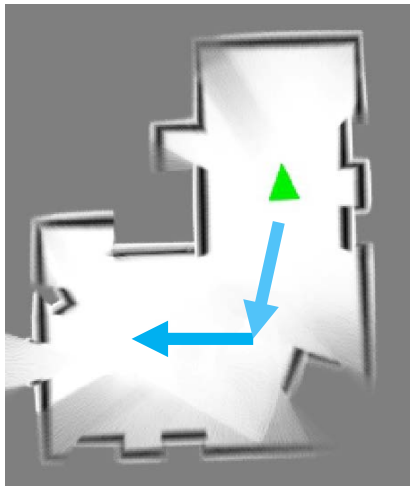

Reinforcement learning

Problem formulation

Agent

Environment

Actions

Problem formulation

Agent

Environment

Actions

How to decide which actions to take?

Policy

Mathematical formulation

Agent

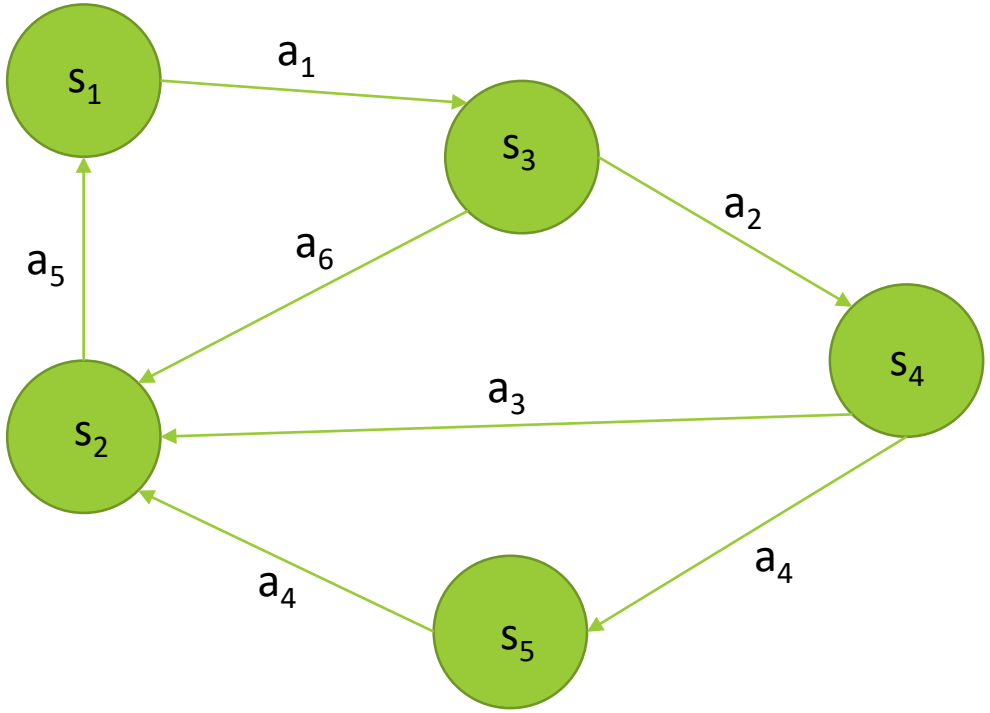

States: s_1, s_2, s_3, s_4,…
Actions: a_1, a_2, a_3, …
Rewards: $r(s_n,a_m)$

Mathematical formulation

Agent

Policy: Action that determines the next state.

Reward: Measure of the outcome of taking an action.

Mathematical formulation

Agent

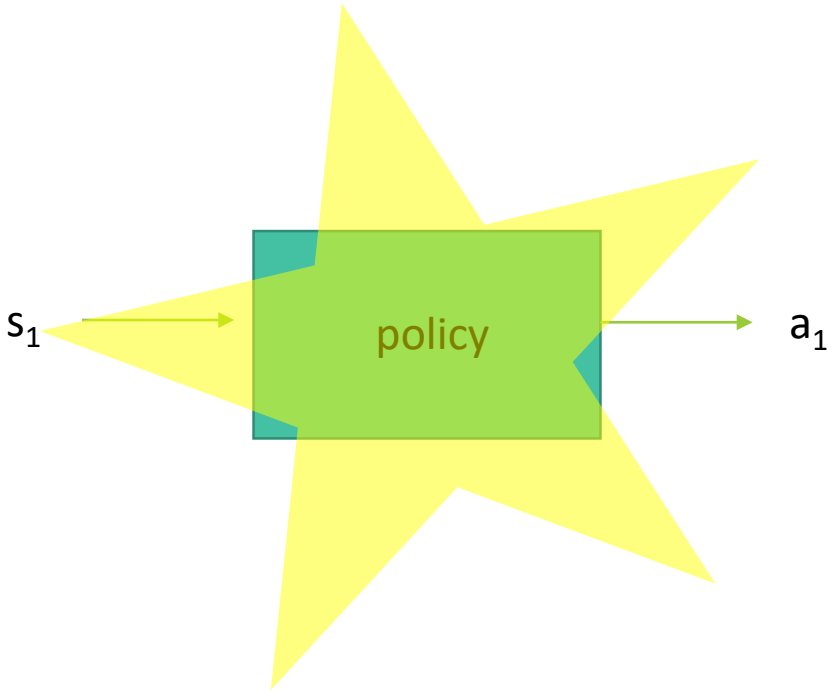

Choosing the ***optimal policy***?

Choosing the optimal policy

Agent

Greedy strategy:

Choosing the optimal policy

Agent

Random strategy:

Multiple random scenarios should be tested before choosing the next action.

Mathematical formulation

Agent

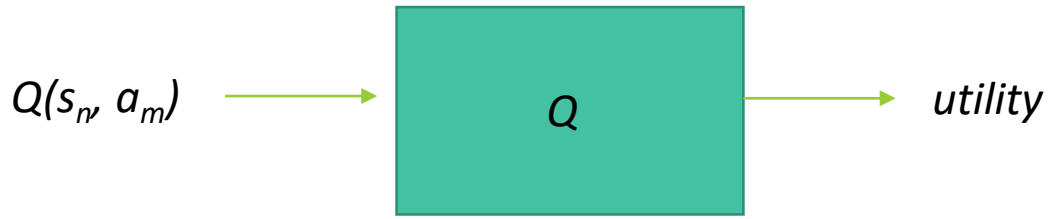

Utility: The long term reward is called a utility.

The utility of performing an action a at *a* state *s* is written as a function *Q(s, a)*, called the utility function.

Mathematical formulation

Agent

If you were given the utility function Q(s,a), how would you use it to derive a policy function?

Mathematical formulation

Agent

If you were given the utility function Q(s,a), how would you use it to derive a policy function?

$$Policy(s) = argmax_a\ Q(s_n, a_m)$$

Mathematical formulation

Agent

$$Q(s,a) = r(s,a) + \gamma\, Q(s',a')$$

Where is the deep part?

Agent

If you were given the utility function Q(s,a), how would you use it to derive a policy function?

$$Policy(s) = argmax_a\ Q(s_n, a_m)$$

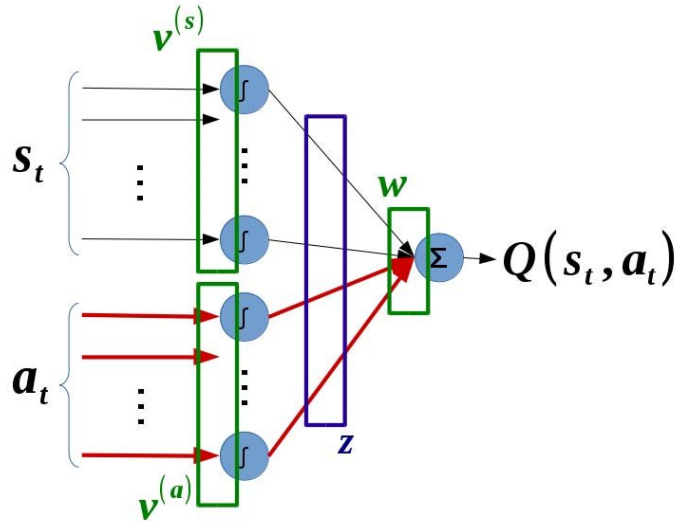

Where is the deep part?

Agent

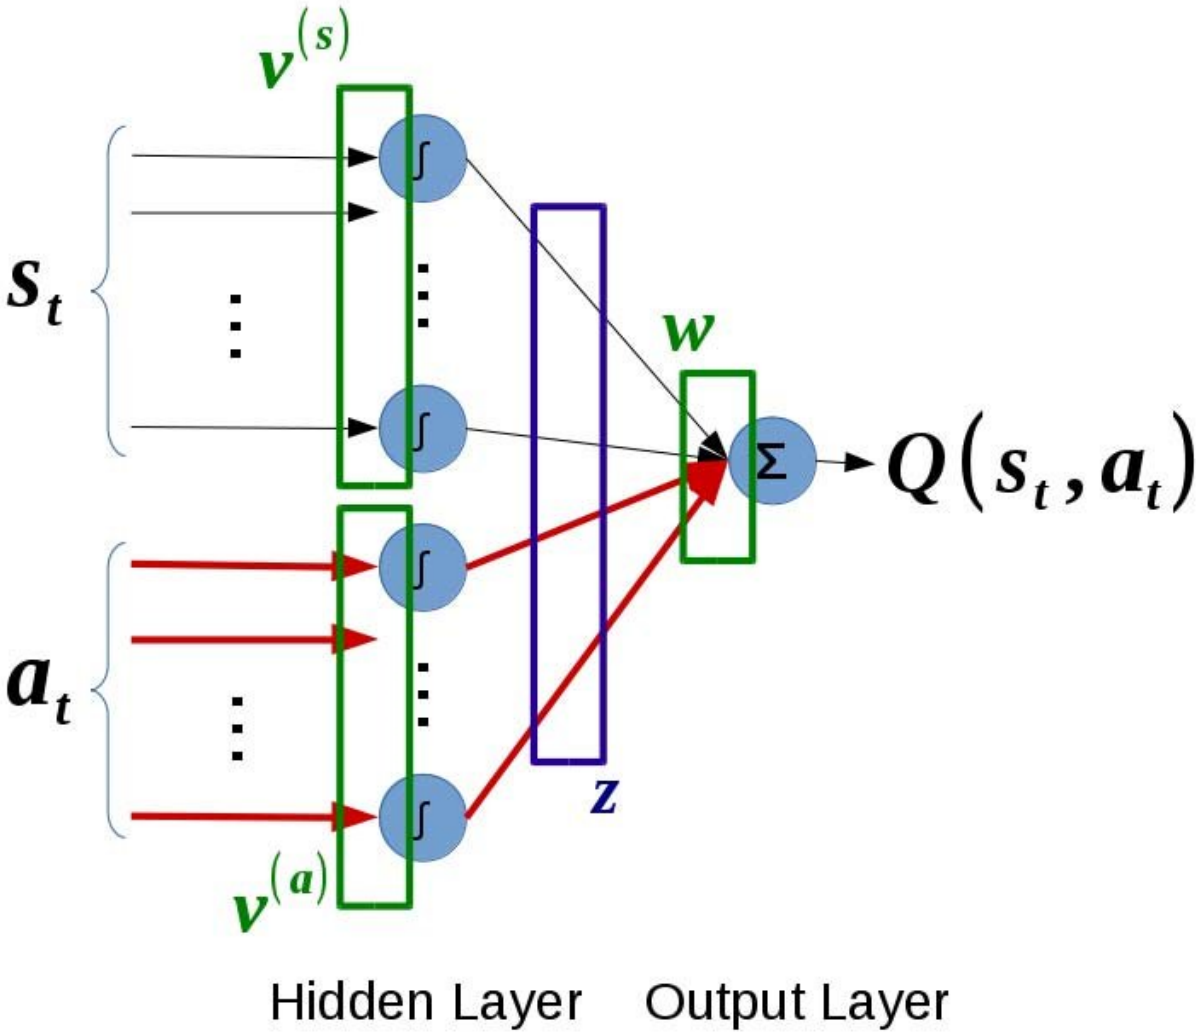

Where is the deep part?

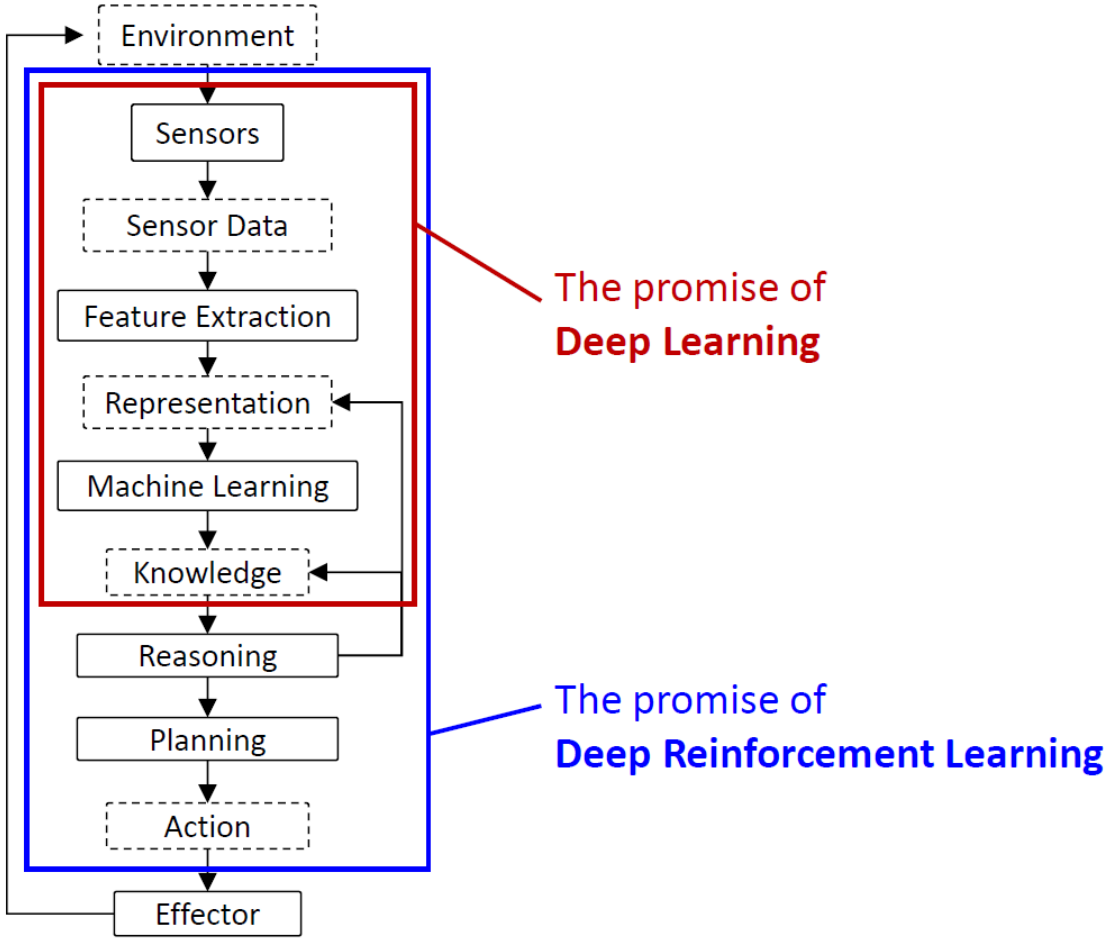

https://deeplearning.mit.edu

Do humans use deep reinforcement learning?

Practice

$ pip install yahoo-finance

Practice

Last 10 days history of the prices

Budget Shares

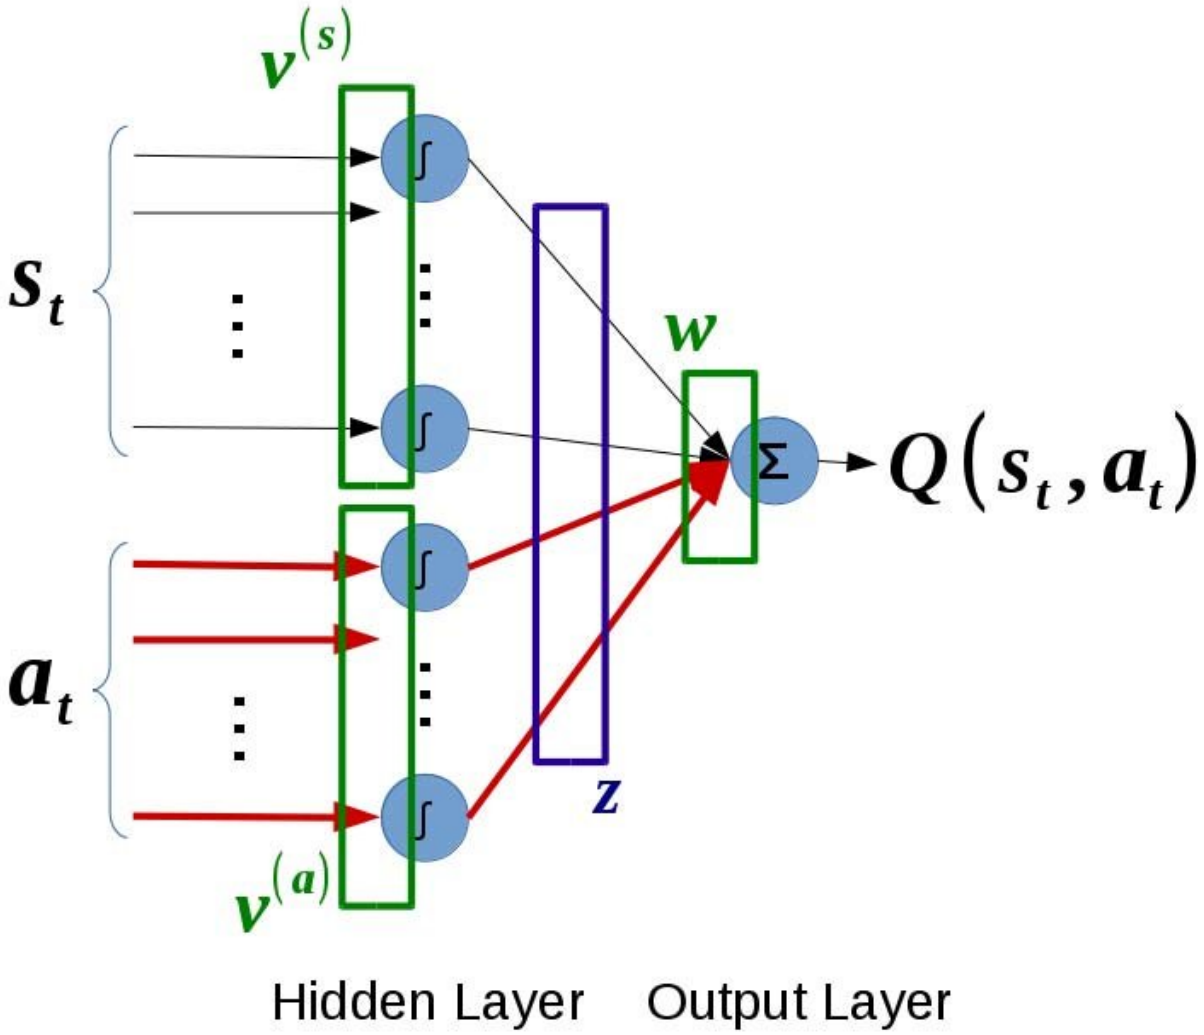

Other applications of reinforcement learning

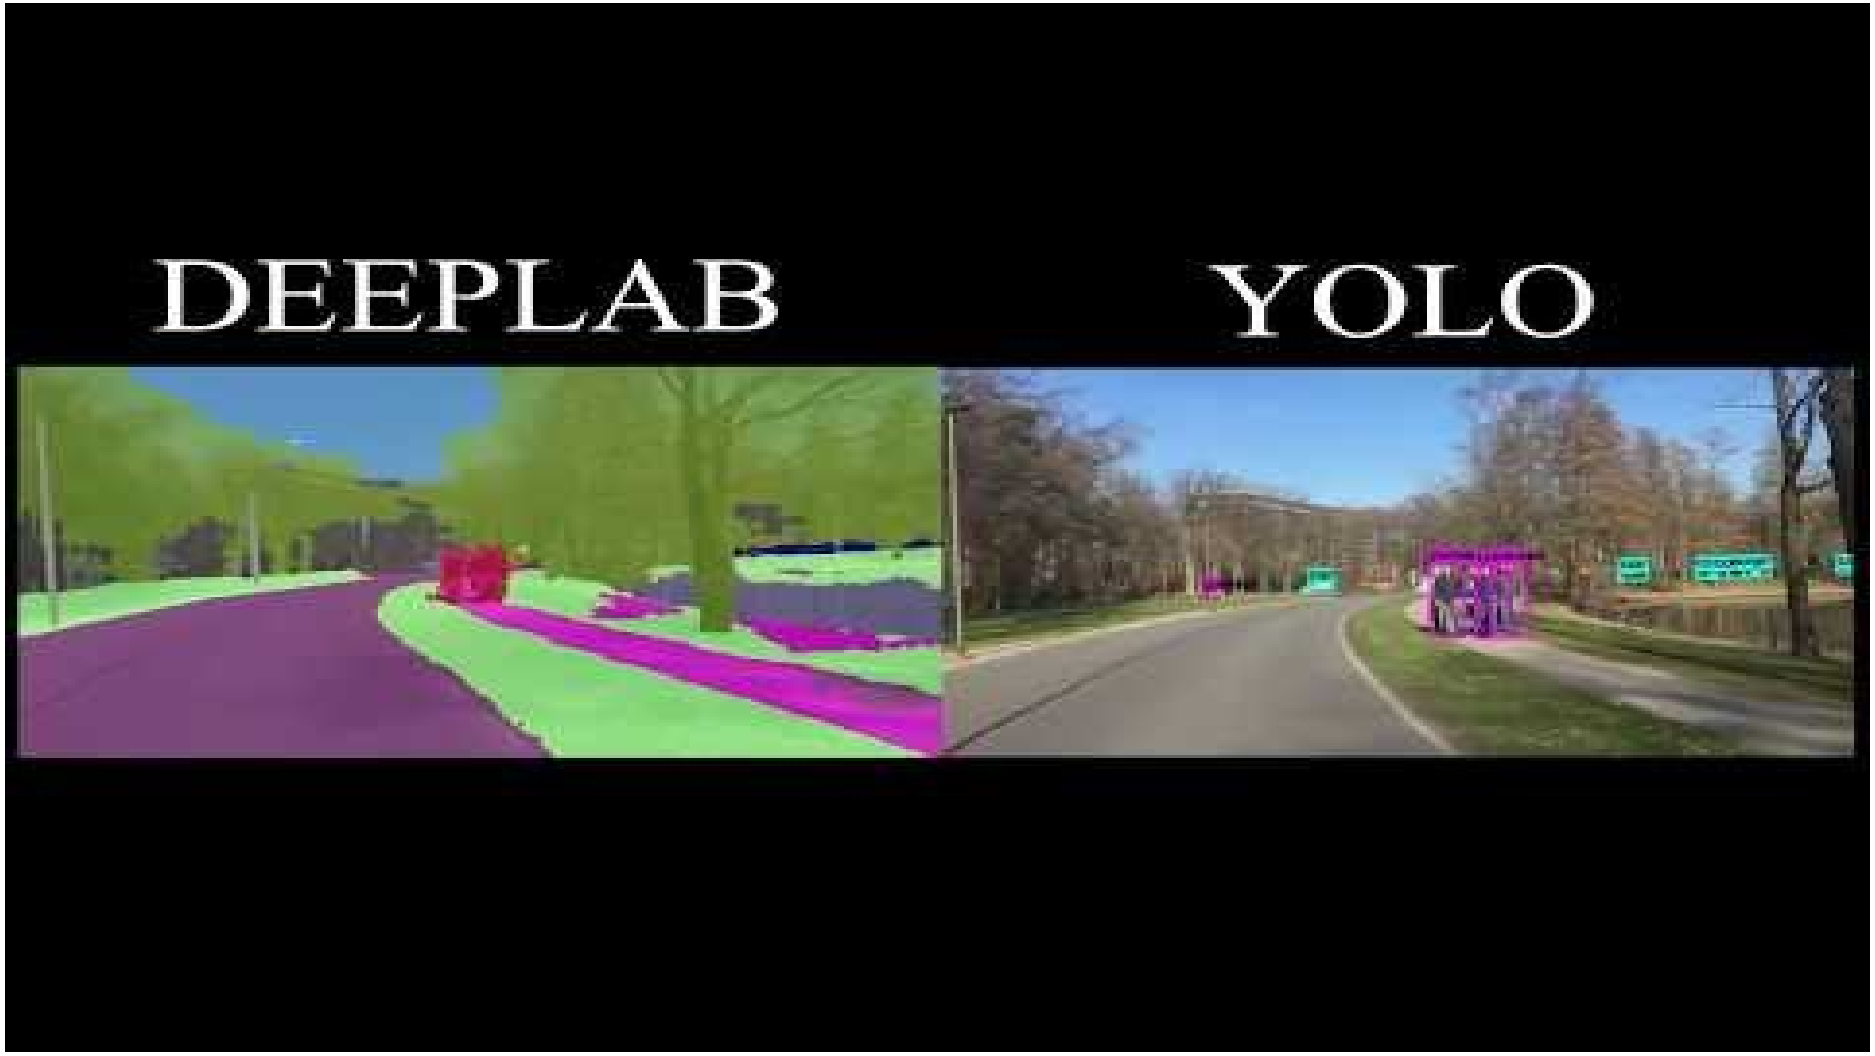

Keras RL

https://github.com/keras-rl/keras-rl

Deep Traffic